

Procuring Performance in Python

Ian J. Bertolacci

Email: ian.bertolacci@gmail.com

Twitter: [@ianbertolacci](https://twitter.com/ianbertolacci)

Github: github.com/ian-bertolacci



About this talk

- General overview of *some* modules and tools that you can use to write more performant code.
- Share!
- I am not an expert in fast Python (yet).
- Code available on github
 - github.com/ian-bertolacci/procuring_python_performace_talk

General Classes of Tools

- Compilers/Interpreters
 - CPython
 - PyPy
- Low-level backed APIs
 - NumPy
- Parallel modules
 - Multiprocessing
 - Threading
- Low-level tie-ins
 - Cython * compiler-y
 - PyCUDA

Why is Python slow?

- Because the implementation is slow.
 - Python's semantics are difficult to provide without a runtime that includes lots of overhead.
- Where does this runtime overhead come from?
 - Dynamic typing
 - More memory requirements
 - Indirect accesses
 - explicit checking
 - Non-contiguous list elements
 - Large integers

Example of Overhead

Python code:

```
a = 1
b = 2
c = a + b
```


This example taken from "Why Python is Slow" by Jake VanderPlas

(<https://jakevdp.github.io/blog/2014/05/09/why-python-is-slow/>)

1. Assign 1 to a
 - a. Set a->PyObject_HEAD->typecode to integer
 - b. Set a->val = 1
2. Assign 2 to b
 - a. Set b->PyObject_HEAD->typecode to integer
 - b. Set b->val = 2
3. call `binary_add(a, b)`
 - a. find typecode in a->PyObject_HEAD
 - b. a is an integer; value is a->val
 - c. find typecode in b->PyObject_HEAD
 - d. b is an integer; value is b->val
 - e. call `binary_add<int, int>(a->val, b->val)`
 - f. result of this is `result`, and is an integer.
4. Create a Python object `c`
 - a. set c->PyObject_HEAD->typecode to integer
 - b. set c->val to `result`

Basic Benchmarks

- Two benchmarks are used here:
 - Fibonacci
 - Both recursive and iterative implementations.
 - $N = 40$
 - Jacobi 2D
 - Stencil computation.
 - 1000^2 grid and 1000 timesteps
 - 5 Giga FLOPS ($1000^2 * 5 * 1000$)
- Machine:
 - Intel i7-6900K @ 3.20GHz
 - 8 cores / 16 hyperthreads
 - Nvidia GeForce GTX 1080
 - 8.5Gb On package RAM
 - 32 Gb RAM




Baseline - Results

- CPython
 - Fibonacci:
 - Iterative:
 - 0.00065 seconds
 - Recursive:
 - 59.2053 seconds
 - Jacobi:
 - 488.244 seconds (~8 minutes)
 - 10.24 Mega FLOPS/second

Baseline - Results

- C
 - Fibonacci
 - Iterative
 - 0.0000003 seconds
 - Recursive
 - 0.433
 - Jacobi
 - Serial
 - 0.897 seconds

OpenMP Jacobi



- Alternative Python interpreter.
 - Probably most popular after CPython
- Written in RPython (a Python derivative). Work your head around that one.
- Pros:
 - Use a just in time (JIT) compiler that compiles python code to lower code closer to machine level
 - No modifications to code required (except...)
- Cons:
 - Limited support for module using CPython's C-API
 - Support getting better, but performance could vary.
 - May require modified libraries

PyPy - Benchmark Results

- Fibonacci
 - Iterative:
 - 0.00007 seconds
 - Recursive:
 - 5.102 seconds
- Jacobi
 - 9.09s (53.6x faster than CPython Jacobi)
 - 549.67 Mega FLOPS/second

PyPy - Related

- There are a gajillion python interpreters and compilers
 - Jython: JVM
 - Pyston: LLVM
 - Pyjion:
 - Hope
 - Falcon
 - PyDron
 - Nuitka: Compiler, almost all of Python.
 - Shed Skin: Compiler, limited subset of Python.
 - Pythran: Compiler, limited subset of Python
 - GT-Py: Intel's interpreter. Adds OpenMP and OpenACC annotations

NumPy

- Non-standard (but wildly popular) module
- Mainly C backed multi-dimensional arrays and some linear algebra tools
- Pros:
 - Powerful array abstractions.
 - Basis of SciPy (scientific python module filled with magic).
 - Mostly written in C/C++/Fortran (fast).
- Cons:
 - Need to modify code to use it.
 - Not a big deal, but would always like to avoid redevelopment.

NumPy - Benchmark Results

- Jacobi
 - 4.592 seconds (106.324x faster than CPython baseline)
 - 1088.850 Mega FLOPS/second

Multiprocessing

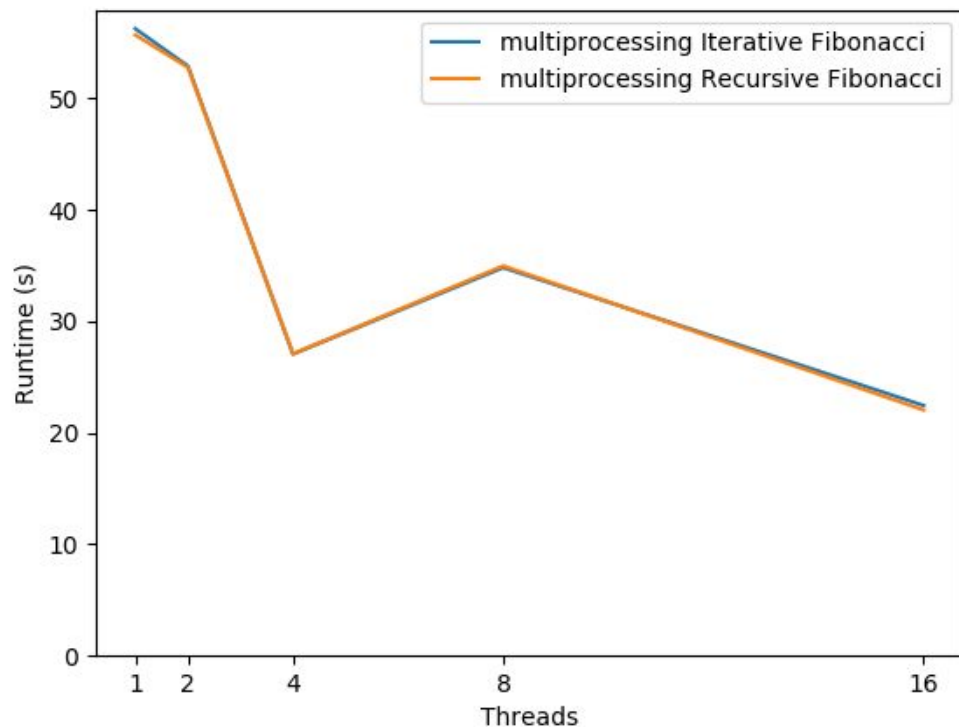
- Standard module.
- Parallelism through processes (not threads).
- Provides mechanisms for usual task-based parallelism.
- Pros:
 - API is fairly standard for a task-based parallelism (create process objects, start, and join them; locks, pipe, semaphore).
 - Provides process pools that can easily map work across them.
 - Supposedly can use multiprocessing on a cluster.
 - Not limited by a Global Interpreter Lock (GIL).
- Cons:
 - Very heavy weight (creates entirely new python interpreter process).
 - Requires you to be quite hands on.

Multiprocessing - Terse Example

```
processes = [  
    Process(  
        target=work_function,  
        args=(work_unit, result_queue)  
    )  
    for work_unit in work_list  
]  
# Start all processes  
for process in processes:  
    process.start()  
# Wait for all processes to stop  
for process in processes:  
    process.join()
```

```
pool = Pool( cpu_count() )  
results = pool.map(  
    work_function,  
    work_list  
)
```

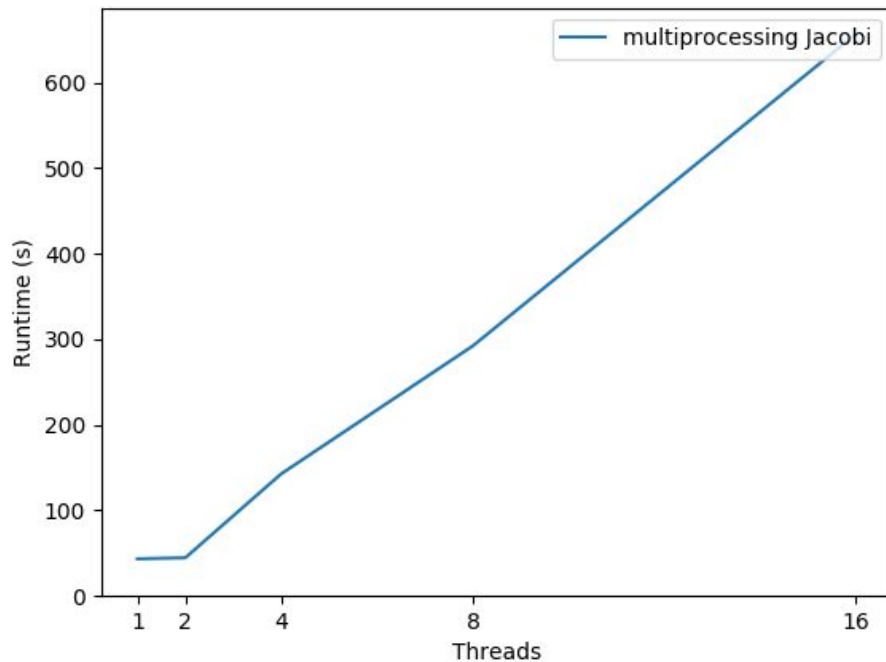
Multiprocessing - Fibonacci Results



Baseline was:

- Iterative: 0.00065 seconds
 - Multiprocessing slower
- Recursive: 59.20 seconds
 - Multiprocessing faster

Multiprocessing - Jacobi Results



Uses NumPy arrays to simplify comms

Baseline was:

- 488.244 seconds
 - Multiprocessing does worse (and keeps getting worse!)

Multiprocessing - Related

- **PyMPI, MPI4Py**
 - Message Passing Interface (MPI) is a very old API for multiprocessing
 - These let you use the API in python
- **Global Arrays**
 - Partitioned Global Addressing Space (PGAS) is a way of thinking about and writing distributed codes.
 - Global Arrays is one such implementation
 - Global Arrays seems to have some Python facing API

Threading

- Standard module
- "Parallelism" via threads
- Very similar to multiprocessing
 - Thread objects, locks, et cetera.
- Pros:
 - API is fairly standard for a task-based parallelism (create process objects, start, and join them; locks, pipe, semaphore)
- Cons:
 - Not concurrent! Bound by GIL.
 - All work and no gain.
- Apparently used for more I/O parallelism...

Threading - Terse Example

```
# Create threads
threads = [
    Thread( target=work_function, args=(work, result_queue) )
    for work in worklist
]
# Start threads
for thread in threads:
    thread.start()
# Join threads
for thread in threads:
    thread.join()
```

Threading - Benchmark Results

- Fibonacci
 - Iterative:
 - 0.0079 seconds
 - Recursive:
 - 176.43 seconds
 -

Threading - Related

- Any kind of asynchronous library is sure to work similarly to the threading module.
 - Probably about a billion of these, primarily used for web services

PyCUDA

- Non-standard module
- Write CUDA kernels in as strings, compile during runtime, and execute.
 - Very similar to how native OpenCL works (and PyOpenCL).
- Pros:
 - Can utilize very powerful hardware in a manner similar (if not identical) to the native API.
 - Spoiler alert: Very fast
- Cons:
 - Single Instruction Multiple Data (SIMD) paradigm limits its application to (essentially) array operations
 - Lots of setup that is confusing, unintuitive, and that can have a performance impact

PyCUDA - Terse Example

```
mod = SourceModule("""
__global__ void multiply_them(float *dest, float *a, float *b, int N){
    const int i = threadIdx.x;
    if( i < N )
        dest[i] = a[i] * b[i];
}
""")
# Compile CUDA function
multiply_them = mod.get_function("multiply_them")
# Create Numpy arrays for input and out
a = numpy.array( data_a ).astype(numpy.float32)
b = numpy.array( data_b ).astype(numpy.float32)
dest = numpy.zeros_like(a)
# Calculate blocksize and grid_size
block_size = 400
grid_size = int( math.ceil(N/float(block_size)) )
# Call CUDA function
multiply_them(
    driver.Out(dest), driver.In(a), driver.In(b), numpy.int32( a.shape[0] ),
    block=(400,1,1), grid=(grid_size), 1)
)
```


PyCUDA - Benchmark Results

- Jacobi
 - 0.0507 seconds (10693x faster CPython, 17.88x faster than C!)
 - 98.619 Giga FLOPS/second

PyCUDA - Related

- PyOpenCL
 - Same but with OpenCL codes
 - Developed by same person/group
- PyChapel
 - Chapel is a high performance programming language from Cray.
 - PyChapel lets you write/use Chapel code and make calls to it from Python

Cython

- Broadly, a Python/Cython-to-C compiler
- Can compile most Python code to a C-API implemented module
- Has Cython language for writing
- Pros:
 - Essentially compiled python
 - Seems to port relatively easily (some minor build process required)
- Cons
 - Does not work in interpreters that don't implement the CPython C-API (limited PyPy support)

Cython - Benchmark Results

- Fibonacci
 - Iterative
 - 8.70227813721e-05s
 - Recursive
 - 15.0813598633s
- Jacobi
 - 221.64 seconds (2.2x faster than CPython)
 - 22.55 Mega FLOPS/second

Cython - Related

- **PyFort**
 - Very similar.
 - Write Python-y Fortran that gets compiled and is callable from Python.
 - Pervasive in SciPy (~25% of project).
- **SWIG**
 - More for creating interfaces between existing code in a language.
- **Grumpy**
 - Python to Go transpiler + runtime
- **Rust?**

Numba

- Annotation system for Python (and maybe an interpreter?)
- Annotate loops with `@jit` or `@vectorize`
 - JIT compiler lowers into LLVM, optimizes, and then to compiles machine code during runtime.
- Pros:
 - Easy to use (on paper). Just annotate
- Cons:
 - Not actually all that easy to use.
 - Confusing type system that does not seem to terminate.
 - When are things arrays? When are they not?
 - Sometimes need to modify code to make work at all.

Numba - Terse Example

```
#Tell numba to JIT foo, infer types
@jit
def foo( a,b ):
    return a+b
```

```
# Tell numba to JIT foo, specifically for these types
@jit([ int32(int32,int32),
      float32(float32,float32),
      float64(float64,float64)])
def bar( a, b ):
    return a-b
```

Numba - Benchmark Results

- Fibonacci
 - Iterative
 - 0.082 seconds (25% slower than CPython baseline)
 - Recursive
 - 1.24985098839 (47x faster than CPython baseline)
- Jacobi
 - 542.145 seconds (11% slower than CPython baseline)

Conclusion

- No magic wand.
 - This is the norm, and is not surprising.
- Quite usable.
 - Writing high performance *always* requires fairly large code modifications.
 - Notably easier to modify in python than with than other languages, like C/C++
- If you need **peak** performance, Python won't be your main application.
 - But Python will work for most people systems
 - You Aren't Google - Ozan Onay
<https://blog.bradfieldcs.com/you-are-not-google-84912cf44afb>
 - Python can be used in other places, such as orchestration of your application(s)
 - Either as a bash replacement
 - Call low level operations (see PyCUDA)